

COSIMA Cookbook: a computational learning module for analysing ocean and sea-ice model output

Navid C. Constantinou^{1,12*}, Julia Neme^{2*}, Wilton Aguiar², Matthis Auger¹⁵, Ashley J. Barnes^{4,9}, Romain Beucher³, Dhruv Bhagtani¹⁰, Christopher Yit Sen Bull⁴, Hannah Dawson¹⁵, Noah Day¹³, Fabio Boeira Dias¹², Edward Doddridge¹⁵, Anupiya Ellepola², Denisse Fierro-Arcos¹⁵, Angus Gibson², Maurice F. Huguenin⁸, Aidan Heerdegen³, Andy McC. Hogg^{3,4}, Ryan M. Holmes⁵, Wilma Huneke², Jemma Jeffree^{2,4}, Andrew E. Kiss², Minghang Li³, Josué Martínez-Moreno¹¹, Jan Jaap Meijer¹⁵, Thomas Moore⁷, Ruth Moorman⁶, Paige Martin³, Adele K. Morrison², James Munroe¹, Aditya Narayanan¹⁴, Micael Oliveira³, Ellie Q. Y. Ong^{4,9}, Max Proft³, Madelaine G. Rosevear^{4,15}, Christina Schmidt¹¹, Taimoor Sohail¹³, Paul Spence^{4,15}, Dougal T. Squire³, Anton Steketee³, Charles Turner³, Felipe Vilela da Silva¹⁵, Marc White^{2,3}, Luwei Yang², Claire Yung², and Jan Zika¹²

¹ 2i2c ² Australian National University, Australia ³ Australian National University, Australia's Climate Simulator (ACCESS-NRI), Australia ⁴ Australian Research Council Centre of Excellence for the Weather of the 21st Century, Australia ⁵ Bureau of Meteorology, Australia ⁶ California Institute of Technology, USA ⁷ Commonwealth Scientific and Industrial Research Organisation, Australia ⁸ GEOMAR Helmholtz Centre for Ocean Research Kiel, Germany ⁹ Monash University, Australia ¹⁰ Princeton University, USA ¹¹ British Antarctic Survey, UK ¹² University of New South Wales Sydney, Australia ¹³ University of Melbourne, Australia ¹⁴ University of Southampton, UK ¹⁵ University of Tasmania, Australia * These authors contributed equally.

DOI: 10.xxxxxx/draft

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

Summary

The COSIMA Cookbook is an open computational learning module for analysing ocean and sea-ice model output in Jupyter notebooks. It has been developed by the Consortium for Ocean-Sea Ice Modelling in Australia (COSIMA; <https://cosima.org.au>) as a community resource for researchers, students, and practitioners working with large gridded datasets, especially output from the ACCESS ocean-sea ice model configurations (e.g. Kiss et al., 2020). The repository combines introductory tutorials with worked analysis examples, all exposed through a browsable Sphinx documentation site (Figure 1) and backed by a citable archived release (Constantinou et al., 2026).

COSIMA Cookbook

Search Go

Navigation

Lessons

Cooking Lessons 101

- Basics
- Advanced

Recipes

Appetisers

Mains

Local Dishes

Resources

Contributing


- Getting started with git
- Working on your recipe
- Tips for making a good recipe
- Submitting a Pull Request
- Reviewing existing Pull Requests

Notebook Guidelines

- Thumbnails
- Images in notebooks

GitHub Repository

Appetisers (easy)



latest

Figure 1: The COSIMA Cookbook documentation website, showing the browsable Sphinx-based landing page used to navigate tutorials and recipes. The live site is available at <https://cosima-recipes.readthedocs.io>.

Following the “Cookbook” concept, the website sections are deliberately named using a gastronomy theme. “Cooking Tutorials” refers to tutorials that teach generic, transferable techniques for handling ocean–sea ice model output. Then comes the main part of any cookbook, its recipes. Here, by “recipes” we mean self-contained notebooks showing how to perform concrete diagnostics and analyses on ocean and sea-ice datasets. First come the “Easy Recipes” that provide a good entry point after the tutorials; then the “Advanced Recipes” that are more elaborate and advanced analysis examples, and last the “Regional Specialties” contains recipes for regional model configurations (Barnes et al., 2024). Together these sections present the documentation as a browsable collection of lessons and recipes gathered into a single cookbook.

Pedagogy and instructional structure are central to the project. The idea is that these recipes provide the starting point for more elaborate analyses. “Cooking Tutorials” introduces generic skills such as loading model output, working with labelled arrays, plotting, and interacting with shared data catalogues. These tutorials lead into domain-focused recipes giving learners an incremental path from first contact with the data ecosystem to adaptation of full workflows for their own science questions.

Within this structure, the tutorials include examples of loading model output through Intake-based catalogues (Intake developers, 2026) and incorporating Cartopy (Elson et al., 2024) for map-based visualisation. The recipe collection also includes a Lagrangian particle-tracking workflow built with Parcels (Delandmeter & Sebille, 2019), as well as analyses that use xarray (Hoyer & Hamman, 2017) and Dask (Dask Development Team, 2016; Rocklin, 2015) to process very large multi-dimensional datasets in parallel, xgcm

(Abernathy et al., 2022) to handle variables defined on staggered finite-volume grids and xESMF (Zhuang et al., 2025) for regridding model output.

The Cookbook grew from a practical need inside the COSIMA community: the tools used to analyse modern ocean model output are powerful, but the gap between package-level documentation and a reproducible end-to-end workflow remains large. Users need to understand not only Python and Jupyter (Kluyver et al., 2016), but also how to navigate high-dimensional model output, shared high-performance computing environments, and domain-specific analysis conventions. The Cookbook addresses that gap by packaging reusable workflows in the same medium in which users actually work.

Statement of Need

Ocean and climate model analysis has a steep entry cost. New users must learn how to find datasets, load them efficiently, interpret metadata, operate on multi-dimensional arrays, and produce scientifically meaningful diagnostics. General-purpose libraries such as xarray (Hoyer & Hamman, 2017) provide the foundations for these tasks, but they do not by themselves show a learner how to translate a research question into a robust analysis workflow for a specific model and computing environment.

That gap is not only pedagogical; it directly affects how quickly model output can be turned into science. For a beginner, becoming productive with a model often requires substantial time just to understand the structure of the output, the relevant conventions, and the computational environment in which analysis is expected to run. There is also substantial scientific labour embedded in many common diagnostics. Some calculations, such as cross-contour transport, require non-trivial development and validation before they can be used confidently. If every new user or project had to reconstruct those workflows independently, a large amount of effort would be repeatedly spent on rebuilding analysis code rather than doing science with the model.

The COSIMA Cookbook fills that need with a domain-specific, openly maintained collection of computational lessons and examples. Its contribution is not a new analysis library; rather, it is a reusable scientific and educational resource that exposes complete, well-documented workflows in the same notebook format in which many researchers actually explore data and communicate methods. By choosing notebooks instead of packaging these examples only as a software library, the project makes intermediate reasoning, methodological choices, and practical execution details visible to learners while still giving more experienced users analysis patterns they can adapt directly for research. This aligns well with JOSE's emphasis on open educational resources that enable computational learning through authentic practice.

Several design decisions make the Cookbook particularly useful for adoption by other groups. First, each notebook is intended to be self-contained and well-documented, so learners can read, run, and modify a complete workflow without reconstructing missing context from scattered notes. Second, the repository distinguishes tutorials from recipes. Tutorials teach transferable skills, while recipes demonstrate how those skills combine in realistic analysis tasks. Third, the project includes contributor guidance and notebook review conventions that encourage new analyses to be written in a reusable and pedagogically clear style. A caveat is that some recipes use high-resolution model output with datasets of order terabytes, so the underlying data are not always practical to distribute or mirror in full. However, beyond the initial data-loading step, which usually occurs early in each recipe, the analysis workflows are generally applicable to most operational ocean-sea ice model outputs that can be loaded with xarray (Hoyer & Hamman, 2017).

Thus, the structure is valuable beyond the immediate COSIMA context. Many research communities maintain model output on shared infrastructure and face the same challenge: turning expert tacit knowledge into examples that newcomers can adapt. The Cookbook

106 offers a model for how a scientific collaboration can capture that knowledge in version-
107 controlled notebooks, publish it as a living resource, and continuously improve it through
108 community contribution.

109 Educational Design and Experience of Use

110 The learning experience is organised as a progression. The documentation points new users
111 first to introductory lessons, including material on loading, slicing, and visualising model
112 output. Learners then move to more advanced tutorials and finally to recipe notebooks
113 that address concrete scientific questions. The categories of “easy” and “advanced” recipes
114 provide a lightweight pedagogical cue about expected complexity and scope, with “regional
115 specialties” specifically covering recipes for regional model configurations (Barnes et al.,
116 2024).

117 The intended mode of use is also explicit. The repository is designed around Jupyter-
118 based analysis on the Australian National Computational Infrastructure, with guidance
119 on running notebooks through ARE/JupyterLab sessions and on accessing shared data
120 holdings. This makes the Cookbook more than a static collection of examples: it is
121 operational documentation for a real analysis environment. At the same time, the
122 notebooks demonstrate broadly transferable patterns for working with labelled geophysical
123 data, so individual lessons can be reused or adapted outside that environment.

124 The project also supports community authorship. Contributors are encouraged to submit
125 new notebooks through pull requests, document their methods clearly, and generalise
126 workflows so that they remain useful to others. In that sense, the Cookbook functions
127 both as a learning module for end users and as a framework for teaching reproducible
128 scientific communication through notebook design and peer review. The easy access to
129 data and analysis provided by the COSIMA Cookbook showcases best practices to the
130 broader oceanographic community and enabled rapid community adoption of the ACCESS
131 ocean and sea-ice model configurations internationally, facilitating to date well over 100
132 peer-reviewed papers and scores of PhD projects.

133 Author order

134 The first two authors contributed equally; authors from the third position onward are
135 listed alphabetically by surname.

136 Acknowledgements

137 We are grateful to the COSIMA community for developing and maintaining the model
138 output, workflows, and pedagogical practices and also for the great community spirit
139 that defines COSIMA. We acknowledge the National Computational Infrastructure (NCI),
140 which is supported by the Commonwealth government of Australia for the computational
141 resources provided under the National Computational Merit Allocation Scheme and the
142 Australian National University Merit Allocation Scheme. We further acknowledge the
143 ACCESS National Research Infrastructure (ACCESS-NRI) ecosystem that make the
144 operational use of these notebooks possible. ACCESS-NRI’s ACCESS-OM2 and ACCESS-
145 OM3 models are enabled by the Australian Government’s National Collaborative Research
146 Infrastructure Strategy (a.k.a. NCRIS). We acknowledge funding from the Australian
147 Research Council under the Centre of Excellence for the Weather of the 21st Century
148 CE230100012, the Linkage Infrastructure, Equipment and Facilities LP160100073 and
149 LP200100406, and the Discovery Project DP240101274.

References

- Abernathy, R. P., Busecke, J. J. M., Smith, T. A., Deauna, J. D., Banihirwe, A., Nicholas, T., Fernandes, F., James, B., Dussin, R., Cherian, D. A., Caneill, R., Sinha, A., Uieda, L., Rath, W., Balwada, D., Constantinou, N. C., Ponte, A., Zhou, Y., Uchida, T., & Thielen, J. (2022). *xgcm* (Version v0.8.1). Zenodo. <https://doi.org/10.5281/zenodo.7348619>
- Barnes, A. J., Constantinou, N. C., Gibson, A. H., Kiss, A. E., Chapman, C., Reilly, J., Bhagtani, D., & Yang, L. (2024). regional-mom6: A Python package for automatic generation of regional configurations for the Modular Ocean Model 6. *Journal of Open Source Software*, 9(100), 6857. <https://doi.org/10.21105/joss.06857>
- Constantinou, N. C., Hogg, A. McC., Gibson, A., Steketee, A., Beucher, R., Proft, M., Heerdegen, A., Neme, J., Holmes, R. M., Morrison, A., Fierro-Arcos, D., Kiss, A. E., Oliveira, M., Yung, C., Doddridge, E., Huneke, W., Bhagtani, D., Ong, E. Q. Y., Munroe, J., ... Narayanan, A. (2026). *COSIMA cookbook: A cookbook of recipes for analysing ocean and sea ice model output*. <https://doi.org/10.5281/zenodo.14353852>
- Dask Development Team. (2016). *Dask: Library for dynamic task scheduling*. <http://dask.pydata.org>
- Delandmeter, P., & Seville, E. van. (2019). The parcels v2.0 lagrangian framework: New field interpolation schemes. *Geoscientific Model Development*, 12(8), 3571–3584. <https://doi.org/10.5194/gmd-12-3571-2019>
- Elson, P., Andrade, E. S. de, Lucas, G., May, R., Hattersley, R., Campbell, E., Comer, R., Dawson, A., Little, B., Raynaud, S., scm72, Snow, A. D., Igolston, Blay, B., Killick, P., lbdreyer, Peglar, P., Wilson, N., Andrew, ... Kirkham, D. (2024). *SciTools/cartopy: v0.24.1* (Version v0.24.1). Zenodo. <https://doi.org/10.5281/zenodo.13905945>
- Hoyer, S., & Hamman, J. J. (2017). xarray: N-D labeled arrays and datasets in Python. *Journal of Open Research Software*, 5(1), 10. <https://doi.org/10.5334/jors.148>
- Intake developers. (2026). *Intake: A lightweight package for finding, investigating, loading and disseminating data*. <https://intake.readthedocs.io/en/reader/>
- Kiss, A. E., Woodward, M. J., Hollings, T., Mu, M., Fiedler, R., Steven, A. D. L., Lenton, A., Matear, R., Chamberlain, M. A., Oke, P. R., Alves, O., Griffies, S. M., Hill, C., Kovach, A., Marsland, S. J., Fiedler, R., Waters, J., Yang, C.-H., Zhou, X., ... Hogg, A. McC. (2020). ACCESS-OM2 v1.0: A global ocean-sea ice model at three resolutions. *Geoscientific Model Development*, 13(2), 401–442. <https://doi.org/10.5194/gmd-13-401-2020>
- Kluyver, T., Ragan-Kelley, B., Pérez, F., Granger, B., Bussonnier, M., Frederic, J., Kelley, K., Hamrick, J., Grout, J., Corlay, S., Ivanov, P., Avila, D., Abdalla, S., & Willing, C. (2016). Jupyter notebooks—a publishing format for reproducible computational workflows. In F. Loizides & B. Schmidt (Eds.), *Positioning and power in academic publishing: Players, agents and agendas* (pp. 87–90). IOS Press. <https://doi.org/10.3233/978-1-61499-649-1-87>
- Rocklin, M. (2015). Dask: Parallel computation with blocked algorithms and task scheduling. In K. Huff & J. Bergstra (Eds.), *Proceedings of the 14th python in science conference* (pp. 130–136). <https://doi.org/10.25080/Majora-7b98e3ed-013>
- Zhuang, J., Dussin, R., Bourgault, P., Huard, D., Banihirwe, A., Raynaud, S., Malevich, B., Schupfner, M., Gauthier, C., Filipe, Levang, S., Almansi, M., Jüling, A., RichardScottOZ, RondeauG, Rasp, S., Smith, T. J., Meyer, A. G., Mares, B., ... Li, X. (2025). *pangeo-data/xESMF: v0.9.2*. <https://doi.org/10.5281/zenodo.4294774>