

COSIMA Cookbook: a community-driven resource for ocean and sea-ice modelling science

Navid C. Constantinou^{3,14*}, Julia Neme^{1*}, Wilton Aguiar¹, Matthis Auger¹³, Ashley J. Barnes^{3,10}, Romain Beucher^{1,2}, Dhruv Bhagtani¹¹, Christopher Yit Sen Bull^{1,2}, Hannah Dawson¹⁷, Noah Day¹⁴, Fabio Boeira Dias¹⁷, Edward W. Doddridge¹⁷, Anupiya Ellepola¹, Denisse Fierro-Arcos¹⁷, Angus Gibson¹, Maurice F. Huguenin⁸, Aidan Heerdegen^{1,2}, Andy McC. Hogg^{1,2,3}, Ryan M. Holmes⁴, Wilma Huneke¹⁰, Jemma Jeffree^{1,3}, Andrew E. Kiss¹, Minghang Li^{1,2}, Josué Martínez-Moreno¹², Jan Jaap Meijer¹⁷, Thomas Moore⁶, Ruth Moorman⁵, Paige Martin^{1,2}, Adele K. Morrison¹, James Munroe⁷, Aditya Narayanan¹⁶, Micael Oliveira^{1,2}, Ellie Q. Y. Ong^{3,10}, Madelaine G. Rosevear^{3,17}, Christina Schmidt¹², Taimoor Sohail¹⁴, Paul Spence^{3,17}, Dougal T. Squire^{1,2}, Anton Steketee^{1,2}, Charles Turner^{1,2}, Felipe Vilela da Silva⁹, Marc White^{1,2}, Luwei Yang¹, Claire Yung¹, and Jan Zika¹⁵

¹ Australian National University, Australia ² Australia's Climate Simulator (ACCESS-NRI), Australia ³ Australian Research Council Centre of Excellence for the Weather of the 21st Century, Australia ⁴ Bureau of Meteorology, Australia ⁵ California Institute of Technology, USA ⁶ Commonwealth Scientific and Industrial Research Organisation, Australia ⁷ Earthdata Associates, Canada ⁸ GEOMAR Helmholtz Centre for Ocean Research Kiel, Germany ⁹ IFREMER - French Institute for Ocean Science, France ¹⁰ Monash University, Australia ¹¹ Princeton University, USA ¹² British Antarctic Survey, UK ¹³ University of Brest, France ¹⁴ University of Melbourne, Australia ¹⁵ University of New South Wales Sydney, Australia ¹⁶ University of Southampton, UK ¹⁷ University of Tasmania, Australia * These authors contributed equally.

DOI: 10.xxxxxx/draft

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

Summary

The COSIMA Cookbook is an open educational resource for analysing ocean and sea-ice model output in Jupyter notebooks (Kluyver et al., 2016). It has been developed by the Consortium for Ocean–Sea Ice Modelling in Australia (COSIMA; <https://cosima.org.au>) as a community resource for researchers, students, and practitioners working with large gridded datasets, especially output from the ACCESS ocean–sea ice model configurations (e.g. Kiss et al., 2020). The Cookbook is developed openly in a GitHub repository and it combines Python Jupyter notebooks of introductory tutorials together with worked analysis examples. All tutorials and examples are exposed through a browsable documentation site built with Sphinx, a documentation generator for Python projects (Komiya et al., 2025); see Figure 1. Tagged versions of the repository are uploaded to a citable archived release (Constantinou et al., 2026).

COSIMA Cookbook

Search Go

Navigation

Lessons

Cooking Tutorials

- Basics
- Advanced

Recipes

Easy Recipes

Advanced Recipes

Regional Specialties

Resources

Contributing


- Getting started with git
- Working on your recipe
- Tips for making a good recipe
- Submitting a Pull Request
- Reviewing existing Pull Requests

Notebook Guidelines

- Thumbnails
- Images in notebooks

GitHub Repository

Easy Recipes



latest

©2025, COSIMA. | Powered by Sphinx 9.1.0 & Alabaster 1.0.0 | [Page source](#)

Figure 1: The COSIMA Cookbook documentation website, showing the browsable Sphinx-based landing page used to navigate tutorials and recipes. The documentation website lives at <https://cosima-recipes.readthedocs.io>.

The COSIMA Cookbook is deliberately organized using a gastronomy theme to help users quickly understand and navigate the website. “Cooking Tutorials” refers to tutorials that teach generic, transferable techniques for handling ocean–sea ice model output. Then comes the main part of any cookbook – its recipes. Here, by “recipes” we mean self-contained notebooks showing how to perform concrete diagnostics and analyses on ocean and sea-ice datasets. First come the “Easy Recipes” that provide a good starting point after the tutorials; then the “Advanced Recipes” that are more elaborate and demonstrate advanced analysis techniques, and lastly the “Regional Specialties” contains recipes for regional model configurations (Barnes et al., 2024). Together these sections present the documentation as a browsable collection of lessons and recipes gathered into a single cookbook. The COSIMA Cookbook is one of several examples of using a gastronomy theme for computational workflows, for example Project Pythia Cookbooks for the general geoscience community and the FAIR Cookbook for working with life science data.

Pedagogy and instructional structure are central to the project. “Cooking Tutorials” introduces generic skills such as loading model output, working with labelled arrays, plotting, and interacting with shared data catalogues. The recipes provide the starting point for more elaborate analyses. These tutorials lead into domain-focused recipes giving learners an incremental path from first contact with the data ecosystem to adaptation of full workflows for their own science questions.

Within this structure, the tutorials leverage and demonstrate open source tools for scientific analysis of large data. This includes examples of loading model output through Intake-based

catalogues (Intake developers, 2026), using Xarray (Hoyer & Hamman, 2017) and Dask (Dask Development Team, 2016; Rocklin, 2015) to process very large multi-dimensional datasets in parallel. The recipe collection includes a Lagrangian particle-tracking workflow built with Parcels (Delandmeter & Seville, 2019), as well as analyses that use xgcm (Abernathy et al., 2022) to handle variables defined on staggered finite-volume grids, and xESMF (Zhuang et al., 2025) for regridding model output. Figures are generated using matplotlib (Hunter, 2007) incorporating Cartopy (Elson et al., 2024) for map-based visualisation.

The Cookbook grew from a practical need inside the COSIMA community (COSIMA Community, 2026): powerful tools for analysing modern ocean model output already existed, but the gap between package-level documentation and reproducible end-to-end workflows left new users largely on their own. Productive use requires fluency not only in Python and Jupyter (Kluyver et al., 2016), but also in navigating high-dimensional, often very large model output, shared high-performance computing environments, and domain-specific analysis conventions. The Cookbook bridges this gap by packaging reusable workflows in the same medium in which researchers actually work.

Statement of Need

Analysing or configuring climate models carries a steep entry cost. New users must learn how to find datasets, load them efficiently, interpret metadata, operate on multi-dimensional arrays, and produce scientifically meaningful diagnostics. General-purpose libraries such as Xarray (Hoyer & Hamman, 2017) provide the foundations for these tasks, but they do not by themselves show a learner how to translate a research question into a robust analysis workflow for a specific model and computing environment. The Cookbook serves as a learning tool for those new to ocean and sea ice science, showing how to interrogate model output and visualise commonly used diagnostics — skills that transfer broadly across models and computing environments.

The gap is not only pedagogical; it directly slows science. Before any analysis can begin, a newcomer typically spends substantial time deciphering the structure of the output, the relevant conventions, and the computational environment. There is also substantial scientific labour embedded in many common diagnostics: some calculations require non-trivial development, understanding of model numerics, and validation before they can be used with confidence — these are the “Advanced Recipes”. If every new user had to reconstruct those workflows independently, enormous effort would be spent rebuilding analysis code rather than doing science.

The COSIMA Cookbook facilitates knowledge sharing and accelerates research with a domain-specific, openly maintained collection of computational lessons and examples. Its contribution is not a new analysis library; rather, it is a reusable scientific and educational resource that exposes complete, well-documented workflows in the same notebook format in which many researchers actually explore data and communicate methods. By choosing notebooks instead of packaging these examples only as a software library, the project makes intermediate reasoning, methodological choices, and practical execution details visible to learners while still giving more experienced users analysis patterns they can adapt directly for research. This aligns well with JOSE’s emphasis on open educational resources that enable computational learning through authentic practice.

Several design decisions make the Cookbook particularly useful for adoption by other groups. First, each notebook is intended to be self-contained and well-documented, so learners can read, run, and modify a complete workflow without reconstructing missing context from scattered notes. Second, the repository distinguishes tutorials from recipes. Tutorials teach transferable skills, while recipes demonstrate how those skills combine in realistic analysis tasks. Third, the project includes contributor guidance and notebook review

conventions that encourage new analyses to be written in a reusable and pedagogically clear style. A caveat is that some recipes use high-resolution model output with datasets of order terabytes, so the underlying data are not always practical to distribute or mirror in full. However, beyond the initial data-loading step, which usually occurs early in each recipe, the analysis workflows are generally applicable to most operational ocean-sea ice model outputs that can be loaded with Xarray (Hoyer & Hamman, 2017).

Thus, the design renders the Cookbook valuable beyond the immediate COSIMA context. Many research communities maintain model output on shared infrastructure and face the same challenge: turning expert tacit knowledge into examples that newcomers can adapt. The Cookbook offers a model for how a scientific collaboration can capture that knowledge in version-controlled notebooks, publish it as a living open resource online, and continuously improve it through community contribution.

The rise of AI-assisted coding tools adds a further dimension to the Cookbook's value. Such tools can lower barriers for newcomers, but they generate far better code when given a well-structured, domain-specific example as context rather than starting from a blank slate. A Cookbook recipe that is self-contained, clearly documented, and validated against real model output is precisely such a starting point. A researcher can hand a recipe together with a description of their own dataset to an AI coding agent and obtain a well-adapted analysis workflow with much less effort than building one from scratch. The Cookbook's relevance thus grows alongside AI adoption: the quality of its examples becomes more, not less, important.

Educational Design and Experience of Use

The learning experience is organised as a progression. The documentation points new users first to introductory lessons, including material on loading, slicing, and visualising model output. Learners then move to more advanced tutorials and finally to recipe notebooks that address concrete scientific questions. The categories of "easy" and "advanced" recipes provide a lightweight pedagogical cue about expected complexity and scope, with "regional specialties" specifically covering recipes for regional model configurations (Barnes et al., 2024).

The intended mode of use is also explicit. The tutorials and recipes showcase Jupyter-based analysis that can be run directly on Australia's National Computational Infrastructure, and include guidance on running notebooks through JupyterLab sessions and on accessing shared data holdings. This makes the Cookbook more than a static collection of examples: it is operational documentation for a real analysis environment. At the same time, the notebooks demonstrate broadly transferable workflows for working with labelled geophysical data, so individual lessons can be reused or adapted for specific research questions or for use on other computational platforms.

The project also supports community authorship. Contributors are encouraged to submit new notebooks through pull requests, document their methods clearly, and generalise workflows so that they remain useful to others. This is particularly valuable as it is relatively rare for the computer code underlying scientific analyses to be peer reviewed with the same care as the manuscript itself. The COSIMA community provides several examples of the recipe ecosystem being extended into full research projects built on peer-authored software repositories on Github, where analysis workflows, methods, and implementation choices are exposed to community scrutiny and reuse. In that sense, the Cookbook functions both as a learning resource for end users and as a framework for teaching reproducible scientific communication through notebook design and peer review. COSIMA workshops and hackathons provide periodic opportunities to review, update, and extend existing recipes, ensuring the Cookbook remains a living resource that keeps pace with evolving tools and community best practices. The easy access to data and analysis

provided by the COSIMA Cookbook has enabled rapid adoption of ACCESS ocean and sea-ice model configurations internationally — not only within physical oceanography but also across adjacent disciplines such as marine ecology, where researchers have used it as an entry point to ACCESS model output. This reach has facilitated well over 100 peer-reviewed papers and more than 20 PhD projects to completion.

Author Order

The first two authors contributed equally; authors from the third position onward are listed alphabetically by surname.

Acknowledgements

This work was supported by computational resources provided by the Australian Government through the National Computational Infrastructure (NCI) under the National Computational Merit Allocation Scheme and the Australian National University Allocation Scheme. We thank the vibrant communities of the Consortium for Ocean–Sea Ice Modelling in Australia (COSIMA; cosima.org.au) and Australia’s climate simulator (ACCESS-NRI; access-nri.org.au) for making the ACCESS-OM2 and ACCESS-OM3 outputs and analysis tools available through the NCI. We further acknowledge funding from the Australian Research Council under the Centre of Excellence for the Weather of the 21st Century CE230100012, the Linkage Infrastructure, Equipment and Facilities LP160100073 and LP200100406, and the Discovery Project DP240101274.

References

- Abernathy, R. P., Busecke, J. J. M., Smith, T. A., Deauna, J. D., Banihirwe, A., Nicholas, T., Fernandes, F., James, B., Dussin, R., Cherian, D. A., Caneill, R., Sinha, A., Uieda, L., Rath, W., Balwada, D., Constantinou, N. C., Ponte, A., Zhou, Y., Uchida, T., & Thielen, J. (2022). *xgcm* (Version v0.8.1). Zenodo. <https://doi.org/10.5281/zenodo.7348619>
- Barnes, A. J., Constantinou, N. C., Gibson, A. H., Kiss, A. E., Chapman, C., Reilly, J., Bhagtani, D., & Yang, L. (2024). regional-mom6: A Python package for automatic generation of regional configurations for the Modular Ocean Model 6. *Journal of Open Source Software*, 9(100), 6857. <https://doi.org/10.21105/joss.06857>
- Constantinou, N. C., Hogg, A. McC., Gibson, A., Steketee, A., Beucher, R., Proft, M., Heerdegen, A., Neme, J., Holmes, R. M., Morrison, A., Fierro-Arcos, D., Kiss, A. E., Oliveira, M., Yung, C., Doddridge, E., Huneke, W., Bhagtani, D., Ong, E. Q. Y., Munroe, J., ... Narayanan, A. (2026). *COSIMA Cookbook: A cookbook of recipes for analysing ocean and sea ice model output*. <https://doi.org/10.5281/zenodo.14353852>
- COSIMA Community. (2026). *COSIMA: The Movie*. <https://youtu.be/bmUYRoPKtIA>.
- Dask Development Team. (2016). *Dask: Library for dynamic task scheduling*. <http://dask.pydata.org>
- Delandmeter, P., & Seville, E. van. (2019). The parcels v2.0 lagrangian framework: New field interpolation schemes. *Geoscientific Model Development*, 12(8), 3571–3584. <https://doi.org/10.5194/gmd-12-3571-2019>
- Elson, P., Andrade, E. S. de, Lucas, G., May, R., Hattersley, R., Campbell, E., Comer, R., Dawson, A., Little, B., Raynaud, S., scm72, Snow, A. D., Iglston, Blay, B., Killick, P., Ildreyer, Peglar, P., Wilson, N., Andrew, ... Kirkham, D. (2024). *SciTools/cartopy: v0.24.1* (Version v0.24.1). Zenodo. <https://doi.org/10.5281/zenodo.13905945>

- 202 Hoyer, S., & Hamman, J. J. (2017). xarray: N-D labeled arrays and datasets in Python.
203 *Journal of Open Research Software*, 5(1), 10. <https://doi.org/10.5334/jors.148>
- 204 Hunter, J. D. (2007). Matplotlib: A 2D graphics environment. *Computing in Science &
205 Engineering*, 9(3), 90–95. <https://doi.org/10.1109/MCSE.2007.55>
- 206 Intake developers. (2026). *Intake: A lightweight package for finding, investigating, loading
207 and disseminating data*. <https://intake.readthedocs.io/en/reader/>
- 208 Kiss, A. E., Woodward, M. J., Hollings, T., Mu, M., Fiedler, R., Steven, A. D. L., Lenton,
209 A., Matear, R., Chamberlain, M. A., Oke, P. R., Alves, O., Griffies, S. M., Hill, C.,
210 Kovach, A., Marsland, S. J., Fiedler, R., Waters, J., Yang, C.-H., Zhou, X., ... Hogg, A.
211 McC. (2020). ACCESS-OM2 v1.0: A global ocean-sea ice model at three resolutions.
212 *Geoscientific Model Development*, 13(2), 401–442. [https://doi.org/10.5194/gmd-13-
213 401-2020](https://doi.org/10.5194/gmd-13-401-2020)
- 214 Kluyver, T., Ragan-Kelley, B., Pérez, F., Granger, B., Bussonnier, M., Frederic, J.,
215 Kelley, K., Hamrick, J., Grout, J., Corlay, S., Ivanov, P., Avila, D., Abdalla, S.,
216 & Willing, C. (2016). Jupyter notebooks—a publishing format for reproducible
217 computational workflows. In F. Loizides & B. Schmidt (Eds.), *Positioning and power
218 in academic publishing: Players, agents and agendas* (pp. 87–90). IOS Press. <https://doi.org/10.3233/978-1-61499-649-1-87>
- 220 Komiya, T., Brandl, G., Turner, A., Bakosi, J.-F., Shimizukawa, T., Andersen, J. L.,
221 Neuhäuser, D., Finucane, S., Lehmann, R., Mason, J., Kampik, T., Addison, J., Magin,
222 J., Dufresne, J., Waltman, J., Eades, D., Cano Rodríguez, J. L., Ronacher, A., Geier,
223 M., ... Freitag, F. (2025). *sphinx-doc/sphinx: Sphinx 9.1.0* (Version v9.1.0). Zenodo.
224 <https://doi.org/10.5281/zenodo.18109073>
- 225 Rocklin, M. (2015). Dask: Parallel computation with blocked algorithms and task
226 scheduling. In K. Huff & J. Bergstra (Eds.), *Proceedings of the 14th python in science
227 conference* (pp. 130–136). <https://doi.org/10.25080/Majora-7b98e3ed-013>
- 228 Zhuang, J., Dussin, R., Bourgault, P., Huard, D., Banihirwe, A., Raynaud, S., Malevich,
229 B., Schupfner, M., Gauthier, C., Filipe, Levang, S., Almansi, M., Jüling, A., Scott,
230 R., RondeauG, Rasp, S., Smith, T. J., Meyer, A. G., Mares, B., ... Li, X. (2025).
231 *pangeo-data/xESMF: v0.9.2*. <https://doi.org/10.5281/zenodo.4294774>